

Launch an AWS Deep Learning AMI

with Amazon EC2

In this step-by-step tutorial, you'll learn how to launch an [AWS Deep Learning AMI](#). The AMIs are machine images loaded with deep learning frameworks that make it simple to get started with deep learning in minutes.

Using the AMI, you can train custom models, experiment with new algorithms, and learn new deep learning skills and techniques. The AMIs come with pre-installed open source deep learning frameworks including TensorFlow, Apache MXNet, PyTorch, Chainer, Microsoft Cognitive Toolkit, Caffe, Caffe2, Theano, and Keras, optimized for high performance on Amazon EC2 instances. The AMIs also offer GPU and CPU-acceleration through pre-configured drivers, and come with popular Python packages.

In the next few minutes, you will launch an EC2 instance using a Deep Learning AMI, connect to the instance via SSH, and access a Jupyter Notebook from your workstation.

The cost of doing this tutorial is the charge for the underlying EC2 instance. We recommend using an EC2 C5.large instance which will generate a charge of less than \$0.13 per hour until you terminate it. This tutorial will instruct you how to terminate the instance to avoid unnecessary charges.

This tutorial requires an AWS account

Step 1: Open the EC2 Console

Sign into the [AWS Management Console](#) with your user name and password to get started. Then type *EC2* in the search bar and open the EC2 service console.

The screenshot shows the AWS Management Console interface. At the top, there's a navigation bar with the AWS logo, 'Services', 'Resource Groups', and a search icon. Below this, a search bar contains the text 'EC2'. A dropdown menu is open, showing search results for 'EC2'. The first result is 'EC2' with the description 'Virtual Servers in the Cloud', which is highlighted with a red box. Other results include 'EFS' (Managed File Storage for EC2) and 'Elastic Container Service' (Run and Manage Docker Containers). Below the search results, there are several categories of services: 'Storage' (S3, EFS, Glacier, Storage Gateway), 'Database' (Relational Database Service, DynamoDB), 'Media Services' (Elastic Transcoder, Kinesis Video Streams, MediaConvert, MediaLive), 'CloudWatch', 'AWS Auto Scaling', 'CloudFormation', 'CloudTrail', 'Config', 'OpsWorks', 'Service Catalog', 'Systems Manager', 'Trusted Advisor', 'Managed Services', 'Mobile Hub', 'AWS AppSync', 'Device Farm', 'Mobile Analytics', 'AR & VR' (Amazon Sumerian), and 'Application Integration' (Step Functions, Amazon MQ, Simple Notification Service, Simple Queue Service, SWF). On the right side, there are 'Helpful tips' and 'Explore AWS' sections. The 'Helpful tips' section includes 'Manage your costs' and 'Create an organization'. The 'Explore AWS' section includes 'Amazon Relational Database Service (RDS)' and 'Real-Time Analytics with Amazon Kinesis'.

Step 2: Configure your instance

In this step, you will set up a server instance with a machine image for deep learning.

Step 2a: Choose the **Launch Instance** button

The screenshot displays the AWS Management Console interface. On the left, a navigation sidebar lists various services, with 'INSTANCES' expanded. The main content area shows the 'Create Instance' page. At the top, it states 'You are using the following Amazon EC2 resources in the US East (N. Virginia) region:' followed by a list of resources: 0 Running Instances, 0 Elastic IPs, 0 Dedicated Hosts, 0 Snapshots, 0 Volumes, 0 Load Balancers, 0 Key Pairs, 1 Security Groups, and 0 Placement Groups. Below this, a blue banner encourages learning more about AWS re:Invent 2017. The 'Create Instance' section includes a 'Launch Instance' button, which is highlighted with a red rectangular box. Below the button, a note specifies that instances will launch in the US East (N. Virginia) region. The 'Service Health' section shows that the US East (N. Virginia) service is operating normally. The 'Availability Zone Status' section shows that the us-east-1a availability zone is operating normally. On the right side of the console, there are links for 'Supported Platforms', 'Additional Information', and 'AWS Marketplace'.

Events

Tags

Reports

Limits

INSTANCES

Instances

Launch Templates

Spot Requests

Reserved Instances

Dedicated Hosts

Scheduled Instances

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Volumes

Snapshots

NETWORK & SECURITY

Security Groups

Elastic IPs

Placement Groups

You are using the following Amazon EC2 resources in the US East (N. Virginia) region:

- 0 Running Instances
- 0 Elastic IPs
- 0 Dedicated Hosts
- 0 Snapshots
- 0 Volumes
- 0 Load Balancers
- 0 Key Pairs
- 1 Security Groups
- 0 Placement Groups

Learn more about the latest in AWS Compute from AWS re:Invent 2017 by viewing the [EC2 Videos](#).

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

Launch Instance

Note: Your instances will launch in the US East (N. Virginia) region

Service Health

Service Status:

- US East (N. Virginia): This service is operating normally

Availability Zone Status:

- us-east-1a: Availability zone is operating normally

Scheduled Events

US East (N. Virginia):

- No events

Supported Platforms

VPC

Default VPC

vpc-ba6d78dc

Resource ID length management

Additional Information

- [Getting Started Guide](#)
- [Documentation](#)
- [All EC2 Resources](#)
- [Forums](#)
- [Pricing](#)
- [Contact Us](#)

AWS Marketplace

Find free software trial products in the AWS Marketplace from the [EC2 Launch Wizard](#). Or try these popular AMIs:

[Barracuda NextGen Firewall F-Series - PAYG](#)

Provided by Barracuda Networks, Inc.

Rating ★★★★★

Starts from \$0.60/hr or from \$4.50/hr

Feedback English (US) © 2008 - 2018 Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy Policy](#) [Terms of Use](#)

Step 2b: Select a AWS Deep Learning AMI

Choose the **AWS Marketplace** tab on the left, and then search for *deep learning ubuntu*. Select the **Deep Learning AMI (Ubuntu)**. You can also select the Base AMI to set up custom builds of deep learning frameworks. Both AMIs are available for Ubuntu or Amazon Linux.

Step 1: Choose an Amazon Machine Image (AMI)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. You can select an AMI provided by AWS, our user community, or the AWS Marketplace; or you can select one of your own AMIs.

The screenshot shows the AWS Marketplace search results for 'deep learning ubuntu'. The search bar at the top contains the text 'deep learning ubuntu'. On the left sidebar, the 'AWS Marketplace' tab is selected. Two AMIs are listed:

- Deep Learning AMI (Ubuntu)**: Sold by Amazon Web Services, 5.0 version, Linux/Unix, Ubuntu 16.04 | 64-bit Amazon Machine Image (AMI) | Updated: 2/27/18. It comes with latest, official binaries of popular deep learning frameworks installed in separate virtual environments. Includes Apache MXNet, TensorFlow, PyTorch, Caffe, Caffe2, ...
- Deep Learning Base AMI (Ubuntu)**: Sold by Amazon Web Services, 3.0 version, Linux/Unix, Ubuntu 16.04 | 64-bit Amazon Machine Image (AMI) | Updated: 1/24/18. It comes with a foundational platform of Nvidia CUDA, cuDNN, GPU drivers, and low-level system libraries for deploying your own custom deep learning environment. For example, for ...

Both AMIs have a 'Select' button next to them.

Step 2c: On the details page, choose **Continue**.

The screenshot shows the details page for the 'Deep Learning AMI (Ubuntu)'. The page is divided into several sections:

- Product Details**:
 - Sold by: Amazon Web Services
 - Customer Rating: 5.0 (2)
 - Latest Version: 5.0
 - Base Operating System: Linux/Unix, Ubuntu 16.04
 - Delivery Method: 64-bit Amazon Machine Image (AMI)
 - License Agreement: End User License Agreement
 - On Marketplace Since: 11/14/17
 - AWS Services Required: Amazon EC2, Amazon EBS
- Pricing Details**:

Instance Type	Software	EC2	Total
R3 Eight Extra Large	\$0.00	\$2.66	\$2.66/hr
M3 Extra Large	\$0.00	\$0.266	\$0.266/hr
R4 16 Extra Large	\$0.00	\$4.256	\$4.256/hr
M5 Extra Large	\$0.00	\$0.192	\$0.192/hr
M4 Extra Large	\$0.00	\$0.20	\$0.20/hr
Graphics Two Extra Large	\$0.00	\$0.65	\$0.65/hr
C3 Quadruple Extra Large	\$0.00	\$0.84	\$0.84/hr
H1 2 Extra Large	\$0.00	\$0.55	\$0.55/hr
High I/O Quadruple Extra Large	\$0.00	\$1.248	\$1.248/hr
T2 Large	\$0.00	\$0.093	\$0.093/hr
C4 Double Extra Large	\$0.00	\$0.398	\$0.398/hr
G2 Eight Extra Large	\$0.00	\$2.60	\$2.60/hr
M5 Large	\$0.00	\$0.096	\$0.096/hr
R3 Double Extra Large	\$0.00	\$0.665	\$0.665/hr
C5 Large	\$0.00	\$0.085	\$0.085/hr
X1e 8 Extra Large	\$0.00	\$6.672	\$6.672/hr
M5 Double Extra Large	\$0.00	\$4.60	\$4.60/hr
		\$0.384	\$0.384/hr

At the bottom right, there are 'Cancel' and 'Continue' buttons. The 'Continue' button is highlighted with a red box.

Step 2d: Select an instance type.

Choose an instance type for your deep learning training and deployment needs, and then click **Review and Launch**. Here we have used a c5.large instance, but you can choose from additional instance types including GPU-based P3 instances.

The screenshot shows the AWS Management Console interface for creating a new instance. The top navigation bar includes the AWS logo, 'Services', 'Resource Groups', and a user profile. Below the navigation bar, a progress bar indicates the current step is '2. Choose Instance Type'. The main content area is titled 'Step 2: Choose an Instance Type' and displays a table of available instance types. The 'c5.large' instance type is selected, highlighted with a red box. The 'Review and Launch' button at the bottom right is also highlighted with a red box.

Instance Type	Instance Class	VCpus	Memory (GiB)	Storage	Network	Accelerated Networking	Eligible for Spot
General purpose	m4.xlarge	16	64	EBS only	Yes	High	Yes
General purpose	m4.10xlarge	40	160	EBS only	Yes	10 Gigabit	Yes
General purpose	m4.16xlarge	64	256	EBS only	Yes	25 Gigabit	Yes
Compute optimized	c5.large	2	4	EBS only	Yes	Up to 10 Gigabit	Yes
Compute optimized	c5.xlarge	4	8	EBS only	Yes	Up to 10 Gigabit	Yes
Compute optimized	c5.2xlarge	8	16	EBS only	Yes	Up to 10 Gigabit	Yes
Compute optimized	c5.4xlarge	16	32	EBS only	Yes	Up to 10 Gigabit	Yes
Compute optimized	c5.9xlarge	36	72	EBS only	Yes	10 Gigabit	Yes
Compute optimized	c5.18xlarge	72	144	EBS only	Yes	25 Gigabit	Yes
Compute optimized	c4.large	2	3.75	EBS only	Yes	Moderate	Yes
Compute optimized	c4.xlarge	4	7.5	EBS only	Yes	High	Yes
Compute optimized	c4.2xlarge	8	15	EBS only	Yes	High	Yes
Compute optimized	c4.4xlarge	16	30	EBS only	Yes	High	Yes
Compute optimized	c4.8xlarge	32	60	EBS only	Yes	10 Gigabit	Yes

At the bottom of the page, there are four buttons: 'Cancel', 'Previous', 'Review and Launch' (highlighted with a red box), and 'Next: Configure Instance Details'.

Step 2e: Launch your instance.

Choose **Launch** on the **Review** page.

Step 7: Review Instance Launch

Please review your instance launch details. You can go back to edit changes for each section. Click **Launch** to assign a key pair to your instance and complete the launch process.

⚠ Your instance configuration is not eligible for the free usage tier
To launch an instance that's eligible for the free usage tier, check your AMI selection, instance type, configuration options, or storage devices. Learn more about [free usage tier](#) eligibility and usage restrictions.

[Don't show me this again](#)

[Edit AMI](#)

AMI Details**Deep Learning AMI (Ubuntu)**

Latest versions of deep learning frameworks pre-installed in separate virtual environments: MXNet, TensorFlow, Caffe, Caffe2, PyTorch, Theano, CNTK, Keras

Root Device Type: ebs Virtualization type: hvm

Hourly Software Fees: \$0.00 per hour on p3.2xlarge instance (Additional taxes may apply.)

Software charges will begin once you launch this AMI and continue until you terminate the instance.

By launching this product, you will be subscribed to this software and agree that your use of this software is subject to the pricing terms and the seller's

[End User License Agreement](#)

Instance Type

[Edit instance type](#)

Instance Type	ECUs	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
p3.2xlarge	23.5	8	61	EBS only	Yes	Up to 10 Gigabit

Security Groups

[Edit security groups](#)

[Cancel](#)

[Previous](#)

Launch

Step 2f: Create a private key file by selecting **Create a new key pair, and download it to a safe location. Then launch your instance.**

Step 7: Review Instance Launch

Please review your instance launch details. You can go back to edit changes for each section. Click **Launch** to assign a key pair to your instance and complete the launch process.

⚠ Your instance configuration is not eligible for the free usage tier
To launch an instance that's eligible for the free usage tier, check your AMI selection, instance type, configuration options, or storage devices. Learn more about [free usage tier](#) eligibility and usage restrictions.

AMI Details

Deep Learning AMI (Ubuntu)
Latest versions of deep learning frameworks pre-installed in separate virtual environments: MXNet, TensorFlow, Caffe, Caffe2, PyTorch, Theano, CNTK, Keras
Root Device Type: ebs Virtualization type: hvm

Hourly Software Fees: \$0.00 per hour on p3.2xlarge instance (Additional taxes may apply.)
Software charges will begin once you launch this AMI and continue until you terminate the instance.

By launching this product, you will be subscribed to this software and agree that your use of this software is subject to the pricing terms and the seller's [End User License Agreement](#)

Instance Type

Instance Type	ECUs	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
p3.2xlarge	23.5	8	61	EBS only	Yes	Up to 10 Gigabit

Security Groups

Select an existing key pair or create a new key pair

A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about removing existing key pairs from a public AMI.

Create a new key pair

Key pair name
my_ami_key

Download Key Pair

You have to download the **private key file** (*.pem file) before you can continue. **Store it in a secure and accessible location.** You will not be able to download the file again after it's created.

Launch Instances

Step 2g: Click **View Instance** to see your instance status.

Launch Status

Your instances are now launching
The following instance launches have been initiated: [i-01f852fcd0a29dcb](#) [View launch log](#)

Get notified of estimated charges
Create billing alerts to get an email notification when estimated charges on your AWS bill exceed an amount you define (for example, if you exceed the free usage tier).

How to connect to your instances

Your instances are launching, and it may take a few minutes until they are in the **running** state, when they will be ready for you to use. Usage hours on your new instances will start immediately and continue to accrue until you stop or terminate your instances.

Click **View Instances** to monitor your instances' status. Once your instances are in the **running** state, you can **connect** to them from the Instances screen. [Find out how to connect to your instances.](#)

▼ **Getting started with your software**

To get started with Deep Learning AMI (Ubuntu) [View Usage Instructions](#) To manage your software subscription [Open Your Software on AWS Marketplace](#)

▼ **Here are some helpful resources to get you started**

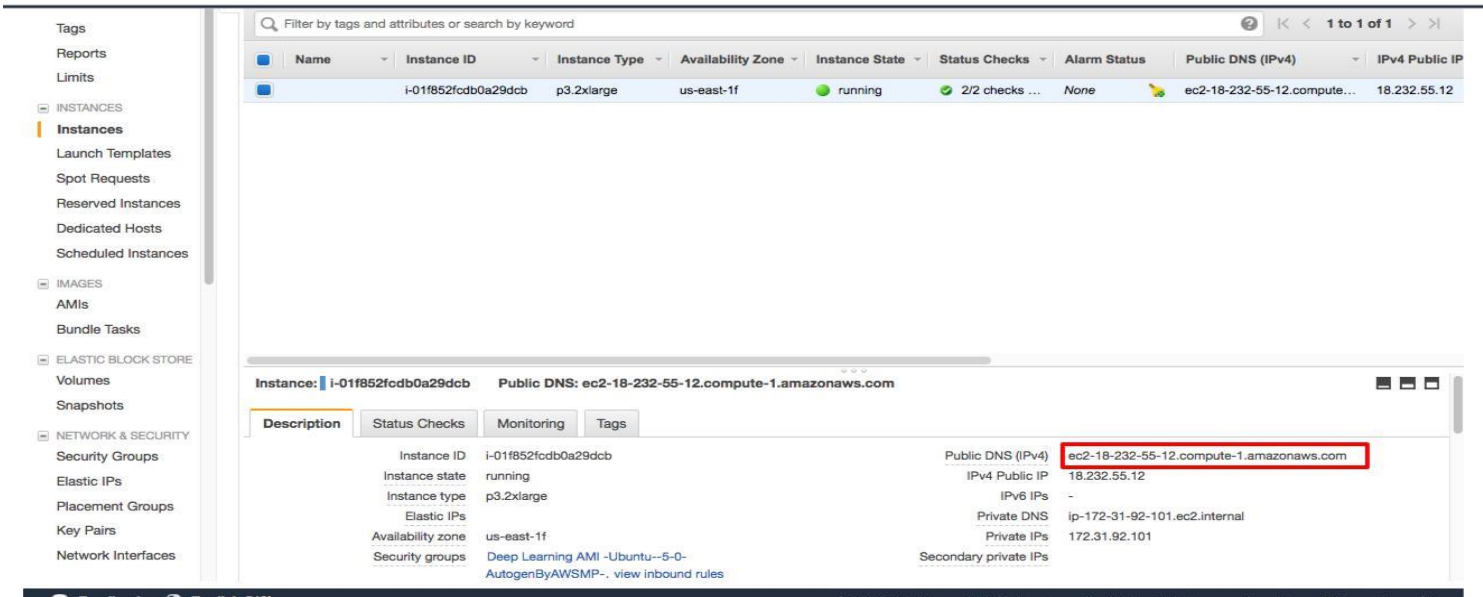
- [How to connect to your Linux instance](#)
- [Amazon EC2: User Guide](#)
- [Learn about AWS Free Usage Tier](#)
- [Amazon EC2: Discussion Forum](#)

While your instances are launching you can also

- [Create status check alarms to be notified when these instances fail status checks. \(Additional charges may apply\)](#)
- [Create and attach additional EBS volumes \(Additional charges may apply\)](#)
- [Manage security groups](#)

[View Instances](#)

Step 2h: Find your instance's public DNS and copy it.



Step 3: Connect to your instance

Here you will use the command line terminal to communicate with the instance on AWS. If you are using Windows, you can use the Command Prompt or download [Git for Windows](#).

Step 3a: Open your command line terminal.

In the terminal, use the following commands to change to the directory where your security key is located, then connect to your instance using SSH. (NOTE: Replace text below in **bold**)

```
cd /Users/your_username/Downloads/
```

```
chmod 0400 <your .pem filename>
```

```
ssh -L localhost:8888:localhost:8888 -i <your .pem filename> ubuntu@<your instance DNS>
```


Step 3b: In the terminal, use the command:

Install tmux

tmux is a terminal multiplexer. It lets you switch easily between several programs in one terminal, detach them (they keep running in the background) and reattach them to a different terminal.

```
sudo apt-get install tmux
```

tmux

```
jupyter notebook --ip=0.0.0.0 --port=8888 --no-browser
```

arguments:

`--ip=<Unicode>` (NotebookApp.ip)

Default: 'localhost'

The IP address the notebook server will listen on.

`--port=<Int>` (NotebookApp.port)

Default: 8888

The port the notebook server will listen on.

`--no-browser`

Don't open the notebook in a browser after startup.

Then copy the URL indicated.

```
JupyterLab v0.27.0
Known labextensions:
[I 20:20:42.654 NotebookApp] Running the core application with no additional extensions or settings
[I 20:20:42.800 NotebookApp] [nb_conda] enabled
/home/ubuntu/anaconda3/lib/python3.6/site-packages/matplotlib/__init__.py:962: UserWarning: Duplicate key in file "/home/ubuntu/.config/matplotlib/matp
c", line #2
(fname, cnt))
/home/ubuntu/anaconda3/lib/python3.6/site-packages/matplotlib/__init__.py:962: UserWarning: Duplicate key in file "/home/ubuntu/.config/matplotlib/matp
c", line #3
(fname, cnt))
[I 20:20:59.434 NotebookApp] sparkmagic extension enabled!
[I 20:20:59.440 NotebookApp] Serving notebooks from local directory: /home/ubuntu
[I 20:20:59.440 NotebookApp] 0 active kernels
[I 20:20:59.440 NotebookApp] The Jupyter Notebook is running at:
[I 20:20:59.440 NotebookApp] http://localhost:8888/?token=5de9f36730db0736ca41d7f8bf5a2da0e16d443086197525
[I 20:20:59.440 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[W 20:20:59.440 NotebookApp] No web browser found: could not locate runnable browser.
[C 20:20:59.440 NotebookApp]

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:
http://localhost:8888/?token=5de9f36730db0736ca41d7f8bf5a2da0e16d443086197525
[I 20:20:59.441 NotebookApp] Starting initial scan of virtual environments...
[I 20:21:53.570 NotebookApp] Found new kernels in environments: conda_theano_p27, conda_caffe_p35, conda_cntk_p27, conda_caffe_p27, conda_mxnet_p27, con
orch_p27, conda_theano_p36, conda_chainer_p27, conda_caffe2_p27, conda_tensorflow_p27, conda_tensorflow_p36, conda_chainer_p36, conda_pytorch_p36, conda
p36, conda_python3, conda_mxnet_p36, conda_python2, conda_anaconda3
```

Step 4: Access your Jupyter Notebook

In this step, you will access your Jupyter Notebook to start using a deep learning framework. If you are connecting to a Jupyter Notebook from a Windows client, you can follow the [steps listed here](https://docs.aws.amazon.com/dlami/latest/devguide/setup-jupyter-configure-client-windows.html). (<https://docs.aws.amazon.com/dlami/latest/devguide/setup-jupyter-configure-client-windows.html>)

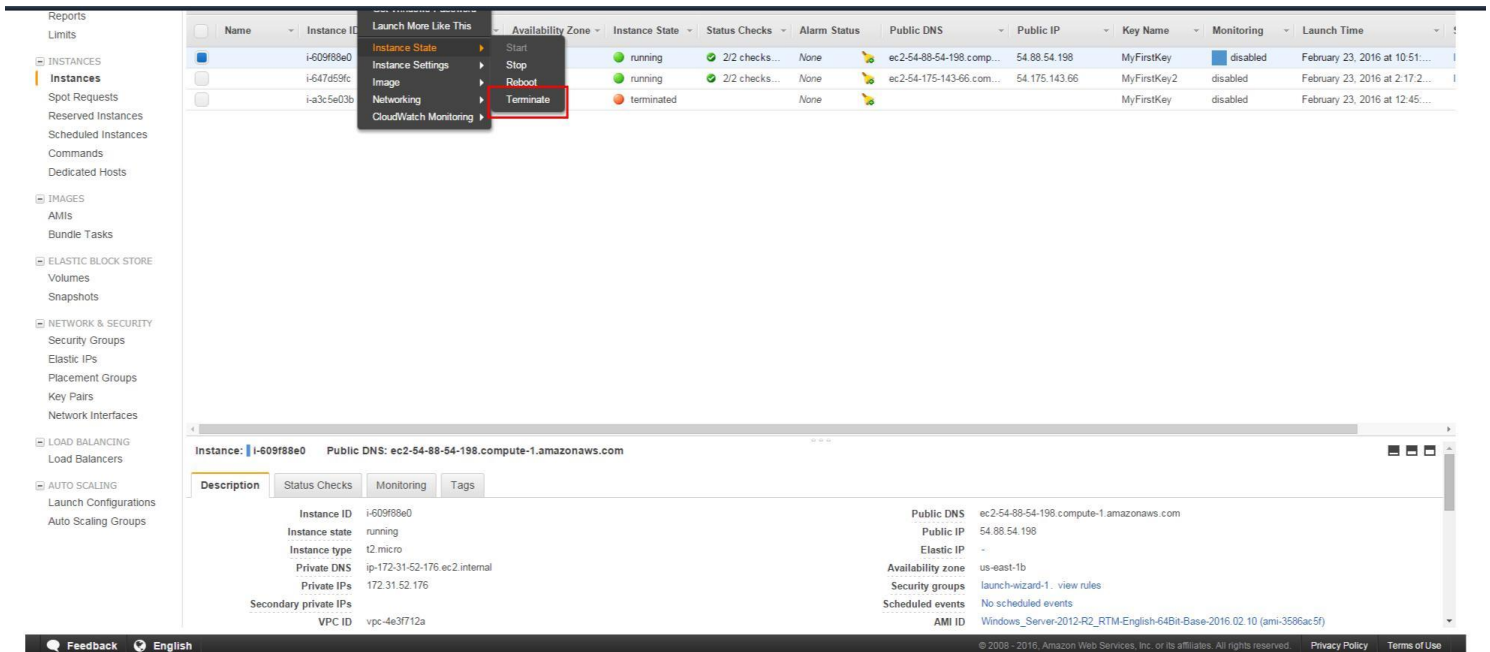
Open a browser window and navigate to the URL indicated in the last step. Choose **New** and start a new notebook using any popular deep learning framework including TensorFlow, Apache MXNet and Gluon, Torch, PyTorch, Caffe, Caffe2, Keras, Theano, Microsoft Cognitive Toolkit, and Chainer.



Step 5: Terminate your instance

Once you're finished, you can easily terminate the instance from the EC2 console. It is a best practice to terminate instances you are no longer using so you don't keep getting charged for them.

Step 5a: Back on the EC2 Console, select the box next to the instance you created. Then click the **Actions** button, navigate to *Instance State*, and click **Terminate**.



The screenshot shows the AWS Management Console interface. On the left, the navigation pane lists various services under categories like INSTANCES, IMAGES, ELASTIC BLOCK STORE, NETWORK & SECURITY, LOAD BALANCING, and AUTO SCALING. The 'INSTANCES' section is expanded, showing a list of instances. The instance 'i-609f88e0' is selected. A context menu is open over this instance, showing options like 'Launch More Like This', 'Instance State', 'Instance Settings', 'Image', 'Networking', and 'CloudWatch Monitoring'. The 'Instance State' sub-menu is open, showing 'Start', 'Stop', 'Reboot', and 'Terminate'. The 'Terminate' option is highlighted with a red box.

Name	Instance ID	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP	Key Name	Monitoring	Launch Time
	i-609f88e0	us-east-1b	running	2/2 checks...	None	ec2-54-88-54-198.compute-1.amazonaws.com	54.88.54.198	MyFirstKey	disabled	February 23, 2016 at 10:51:...
	i-647d59fc	us-east-1b	running	2/2 checks...	None	ec2-54-175-143-66.compute-1.amazonaws.com	54.175.143.66	MyFirstKey2	disabled	February 23, 2016 at 2:17:2...
	i-a3c5e03b	us-east-1b	terminated		None			MyFirstKey	disabled	February 23, 2016 at 12:45:...

Below the table, the details for instance 'i-609f88e0' are shown. The 'Description' tab is active, displaying the following information:

- Instance ID: i-609f88e0
- Instance state: running
- Instance type: t2.micro
- Private DNS: ip-172-31-52-176.ec2.internal
- Private IPs: 172.31.52.176
- Secondary private IPs: (none)
- VPC ID: vpc-4e3f712a
- Public DNS: ec2-54-88-54-198.compute-1.amazonaws.com
- Public IP: 54.88.54.198
- Elastic IP: (none)
- Availability zone: us-east-1b
- Security groups: launch-wizard-1 (view rules)
- Scheduled events: No scheduled events
- AMI ID: Windows_Server-2012-R2_RTM-English-64Bit-Base-2016.02.10 (ami-3586ac5f)

Step 5b: You will be asked to confirm your termination - **Yes, Terminate.**

Note: This process can take several seconds to complete. Once your instance has been terminated, the Instance State will change to *terminated* on your EC2 Console.

The screenshot shows the AWS Management Console interface. On the left, there is a navigation menu with categories like Reports, Limits, INSTANCES, SPOT REQUESTS, ELASTIC BLOCK STORE, NETWORK & SECURITY, LOAD BALANCING, and AUTO SCALING. The main area displays a table of EC2 instances. Two instances are listed: i-609f88e0 (t2.micro, us-east-1b, running) and i-647d59fc (t2.micro, us-east-1d, running). A third instance, i-a3c5e03b, is shown as terminated. A 'Terminate Instances' dialog box is open in the center, displaying a warning about EBS volume deletion and asking for confirmation to terminate the instance i-609f88e0. The 'Yes, Terminate' button is highlighted with a red box.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS	Public IP	Key Name	Monitoring	Launch Time
	i-609f88e0	t2.micro	us-east-1b	running	2/2 checks...	None	ec2-54-88-54-198.compute-1.amazonaws.com	54.88.54.198	MyFirstKey	disabled	February 23, 2016 at 10:51:...
	i-647d59fc	t2.micro	us-east-1d	running	2/2 checks...	None	ec2-54-175-143-66.compute-1.amazonaws.com	54.175.143.66	MyFirstKey2	disabled	February 23, 2016 at 2:17:2...
	i-a3c5e03b	t2.micro	us-east-1d	terminated		None			MyFirstKey	disabled	February 23, 2016 at 12:45:...

Terminate Instances

Warning
On an EBS-backed instance, the default action is for the root EBS volume to be deleted when the instance is terminated. Storage on any local drives will be lost.

Are you sure you want to terminate these instances?

- i-609f88e0 (ec2-54-88-54-198.compute-1.amazonaws.com)

Cancel **Yes, Terminate**

Instance: **i-609f88e0** Public DNS: ec2-54-88-54-198.compute-1.amazonaws.com

Description	Status Checks	Monitoring	Tags
Instance ID: i-609f88e0	Instance state: running	Public DNS: ec2-54-88-54-198.compute-1.amazonaws.com	
Instance type: t2.micro		Public IP: 54.88.54.198	
Private DNS: ip-172-31-52-175.ec2.internal		Elastic IP: -	
Private IPs: 172.31.52.176		Availability zone: us-east-1b	
Secondary private IPs: -		Security groups: launch-wizard-1, view rules	
VPC ID: vpc-4e3f712a		Scheduled events: No scheduled events	
		AMI ID: Windows_Server-2012-R2_RTM-English-64Bit-Base-2016.02.10 (ami-3586ac5f)	

Feedback English

© 2008 - 2016, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

Now that you've launched an AWS Deep Learning AMI, you can easily run tutorials for computer vision, natural language processing, recommendation systems, and more using the deep learning framework of your choice. Happy modeling!